

4.9)

Consider the following code, which multiplies two vectors that contain single-precision complex values:

```
for (i=0; i < 300; i++) {  
    c_re[i] = a_re[i] * b_re[i] - a_im[i] * b_im[i];  
    c_im[i] = a_re[i] * b_im[i] + a_im[i] * b_re[i];  
}
```

Assume that the processor runs at 700 MHz and has a maximum vector length of 64. The load/store unit has a start-up overhead of 15 cycles; the multiply unit, 8 cycles; and the add/subtract unit, 5 cycles.

a) What is the arithmetic intensity of this kernel? Justify your answer.

There are four variables being used loaded each loop:

- a_re[i]
- b_re[i]
- a_im[i]
- b_im[i]

There are 2 variables being written in each loop:

- c_re[i]
- c_im[i]

There are 6 operations per loop:

- * (4)
- + (1)
- - (1)

Therefore the arithmetic intensity of the loop is $(4+2) / 6 = 1$

b) Convert this loop into RV64V assembly code using strip mining.

Based on the code shown on p. 296 of Hennessey and Patterson

vsetdcfg	2 DP FP	Enable 2 64b Fl. Pt. registers
vld	v4,x4	load vector c_{re}
vld	v5,x5	load vector c_{im}
setvl (loop)	t0,a0	$vl = t0 = \min(mvl, n)$
slli	t1,t0,3	$t1 = vl * 8$
vld	v0,x0	load vector a_{re}
add	x0,x0,t1	increment pointer to a_{re} by $vl*8$
vld	v1,x1	load vector b_{re}
add	x1,x1,t1	increment pointer to b_{re} by $vl*8$
vld	v2,x2	load vector b_{im}
add	x2,x2,t1	increment pointer to b_{im} by $vl*8$
vmul	v0,v2,v0	$a_{re} * b_{im}$
vld	v3,x3	load vector a_{im}
add	x3,x3,t1	increment pointer to a_{im} by $vl*8$
vmul	v3,v1,v1	$a_{im} * b_{re}$
vmul	v3,v2,v2	$a_{im} * b_{im}$
vmul	v0,v1,v3	$a_{re} * b_{re}$
vadd	v0,v1,v0	$(a_{re} * b_{im}) - (a_{im} * b_{re})$
vsub	v3,v2,v1	$(a_{re} * b_{re}) - (a_{im} * b_{im})$
add	x4,x4,t1	increment pointer to c_{re} by $vl*8$
add	x5,x5,t1	increment pointer to c_{im} by $vl*8$
vst	v1,x4	store the value of $(a_{re} * b_{re}) - (a_{im} * b_{im})$ in c_{re}
vst	v0,x5	store the value of $(a_{re} * b_{im}) - (a_{im} * b_{re})$ in c_{im}
bnez	a0,loop	repeat if $n \neq 0$
vdisable		disable vector registers

- c) Assuming chaining and a single memory pipeline, how many chimes are required? How many clock cycles are required per complex result value, including start-up overhead?

vld	vld	load a_re and b_re
vmul	vld	a_re * b_re and load a_im
vld	vmul	load b_im and a_im*b_im
vsub	vst	do subtraction and store into c_re
vmul	vld	a_re*b_im and load a_re
vmul	vld	a_im*b_re and load b_re
vadd	vst	do addition and store into c_im

Here **6 chimes** are required, with one extra overhead for the initial startup for the loading of a_re and b_re.

Assume a maximum vector length of 64. The load/store unit has a start-up overhead of 15 cycles; the multiply unit, 8 cycles; and the add/subtract unit, 5 cycles.

Therefore total cycles per iteration =

chimes * elements + overhead

Here the overhead is the first loop:

6 loads * 15 cycles + 4 multiplication * 8 cycles + 2 add/sub * 4 cycles = 132

So total cycles per iteration = 6 * 64 + 132 = 516

Cycles per result = 516/(64 * 2 (because 2 vectors are calculated)) = 4