


Fundamentals of Quantitative Design and Analysis

Professor Hugh C. Lauer

Professor Therese M. Smith

CS-4515, System Programming Concepts

(Slides include copyright materials from Computer Architecture: A Quantitative Approach, 5th ed., by Hennessy and Patterson and from Computer Organization and Design, 4th ed. by Patterson and Hennessy)



Chapter 1

Fundamentals of Quantitative Design and Analysis

Computer Technology

■ Performance improvements:

- Improvements in semiconductor technology
 - Feature size, clock speed
- Improvements in computer architectures
 - Enabled by HLL compilers, UNIX
 - Lead to RISC architectures
- Together have enabled:
 - Lightweight computers
 - Productivity-based managed/interpreted programming languages

Single Processor Performance




Figure 1.1

Current Trends in Architecture

- **Cannot continue to leverage Instruction-Level parallelism (ILP)**
 - Single processor performance improvement ended in 2003
- **New models for performance:**
 - Data-level parallelism (DLP)
 - Thread-level parallelism (TLP)
 - Request-level parallelism (RLP)
- **These require explicit restructuring of applications**

Classes of Computers

■ Personal Mobile Device (PMD)

- e.g. smart phones, tablet computers
- Emphasis on energy efficiency and real-time

■ Desktop Computing

- Emphasis on price-performance

■ Servers

- Emphasis on availability, scalability, throughput

■ Clusters / Warehouse Scale Computers

- Used for “Software as a Service (SaaS)”
- Emphasis on availability and price-performance
- Sub-class: Supercomputers, emphasis: floating-point performance and fast internal networks

■ Internet of Things/Embedded Computers

- Emphasis: price

Parallelism

- **Classes of parallelism in applications:**
 - Data-Level Parallelism (DLP)
 - Task-Level Parallelism (TLP)

- **Classes of architectural parallelism:**
 - Instruction-Level Parallelism (ILP)
 - Vector architectures/Graphic Processor Units (GPUs)
 - Thread-Level Parallelism
 - Request-Level Parallelism

Flynn's Taxonomy

- **Single instruction stream, single data stream (SISD)**
- **Single instruction stream, multiple data streams (SIMD)**
 - Vector architectures
 - Multimedia extensions
 - Graphics processor units
- **Multiple instruction streams, single data stream (MISD)**
 - No commercial implementation
- **Multiple instruction streams, multiple data streams (MIMD)**
 - Tightly-coupled MIMD
 - Loosely-coupled MIMD

These are key terms for all Computer Science topics.

Defining Computer Architecture

■ “Traditional” view of computer architecture:–

- Instruction Set Architecture (ISA) design
- i.e. decisions regarding:
 - registers, memory addressing, addressing modes, instruction operands, available operations, control flow instructions, instruction encoding

■ “Real” computer architecture:

- Specific requirements of the target machine
- Design to maximize performance within constraints:– cost, power, and availability
- Includes ISA, microarchitecture, hardware

Instruction Set Architecture (ISA)

■ Classes of ISA

- General-purpose registers
- Register-memory vs load-store

■ RISC-V registers

- 32 g.p., 32 f.p.

Register	Name	Use	Saver
x0	zero	constant 0	n/a
x1	ra	return addr	caller
x2	sp	stack ptr	callee
x3	gp	gbl ptr	
x4	tp	thread ptr	
x5-x7	t0-t2	temporaries	caller
x8	s0/fp	saved/ frame ptr	callee

Register	Name	Use	Saver
x9	s1	saved	callee
x10-x17	a0-a7	arguments	caller
x18-x27	s2-s11	saved	callee
x28-x31	t3-t6	temporaries	caller
f0-f7	ft0-ft7	FP temps	caller
f8-f9	fs0-fs1	FP saved	callee
f10-f17	fa0-fa7	FP arguments	callee
f18-f27	fs2-fs21	FP saved	callee
f28-f31	ft8-ft11	FP temps	caller

Instruction Set Architecture

■ Memory addressing

- RISC-V: byte addressed, aligned accesses faster

■ Addressing modes

- RISC-V: Register, immediate, displacement (base+offset)
- Other examples: autoincrement, indexed, PC-relative

■ Types and size of operands

- RISC-V: 8-bit, 32-bit, 64-bit

Instruction Set Architecture

■ Operations

- RISC-V: data transfer, arithmetic, logical, control, floating point
- See Fig. 1.5 in text

■ Control flow instructions

- Use content of registers (RISC-V) vs. status bits (x86, ARMv7, ARMv8)
- Return address in register (RISC-V, ARMv7, ARMv8) vs. on stack (x86)

■ Encoding

- Fixed (RISC-V, ARMv7/v8 except compact instruction set) vs. variable length (x86)

Trends in Technology

- **Integrated circuit technology (Moore's Law)**
 - Transistor density: 35%/year
 - Die size: 10-20%/year
 - Integration overall: 40-55%/year
- **DRAM capacity: 25-40%/year (slowing)**
 - 8 Gb (2014), 16 Gb (2019), possibly no 32 Gb
- **Flash capacity: 50-60%/year**
 - 8-10X cheaper/bit than DRAM
- **Magnetic disk capacity: recently slowed to 5%/year**
 - Density increases may no longer be possible, maybe increase from 7 to 9 platters
 - 8-10X cheaper/bit than Flash
 - 200-300X cheaper/bit than DRAM

Bandwidth and Latency


■ Bandwidth or throughput

- Total work done in a given time
- 32,000-40,000X improvement for processors
- 300-1200X improvement for memory and disks

■ Latency or response time

- Time between start and completion of an event
- 50-90X improvement for processors
- 6-8X improvement for memory and disks

Bandwidth and Latency



Log-log plot of bandwidth and latency milestones

Figure 1.9

Transistors and Wires

■ Feature size

- Minimum size of transistor or wire in x or y dimension
- 10 microns in 1971 to .011 microns in 2017
- Transistor performance scales linearly
 - Wire delay does not improve with feature size!
- Integration density scales quadratically

Power and Energy

- **Problem: Get power in, get power out**
- **Thermal Design Power (TDP)**
 - Characterizes sustained power consumption
 - Used as target for power supply and cooling system
 - Lower than peak power (1.5X higher), higher than average power consumption
- **Clock rate can be reduced dynamically to limit power consumption**
- **Energy per task is often a better measurement**

Dynamic Energy and Power

■ Dynamic energy

- Transistor switch from 0 -> 1 or 1 -> 0
- $\frac{1}{2} \times \text{Capacitive load} \times \text{Voltage}^2$


■ Dynamic power

- $\frac{1}{2} \times \text{Capacitive load} \times \text{Voltage}^2 \times \text{Frequency switched}$

■ Reducing clock rate reduces power, not energy

Power


- Intel 80386 consumed ~ 2 W
- 3.3 GHz Intel Core i7 consumes 130 W
- Heat must be dissipated from 1.5×1.5 cm chip
- This is the limit of what can be cooled by air



Reducing Power

■ Techniques for reducing power:—


- Do nothing well
- Dynamic Voltage-Frequency Scaling



- Low power state for DRAM, disks
- Overclocking, turning off cores

Static Power

- **Static power consumption**
 - 25-50% of total power
 - $\text{Current}_{\text{static}} \times \text{Voltage}$
 - Scales with number of transistors
 - To reduce: power gating



Trends in Cost

- **Cost driven down by learning curve**
 - Yield
- **DRAM: price closely tracks cost**
- **Microprocessors:— price depends on volume**
 - 10% less for each doubling of volume

Integrated Circuit Cost

■ Integrated circuit

$$\text{Cost of integrated circuit} = \frac{\text{Cost of die} + \text{Cost of testing die} + \text{Cost of packaging and final test}}{\text{Final test yield}}$$

$$\text{Cost of die} = \frac{\text{Cost of wafer}}{\text{Dies per wafer} \times \text{Die yield}}$$

$$\text{Dies per wafer} = \frac{\pi \times (\text{Wafer diameter}/2)^2}{\text{Die area}} - \frac{\pi \times \text{Wafer diameter}}{\sqrt{2} \times \text{Die area}}$$

■ Bose-Einstein formula:

$$\text{Die yield} = \text{Wafer yield} \times 1/(1 + \text{Defects per unit area} \times \text{Die area})^N$$

- Defects per unit area = 0.016-0.057 defects per square cm (2010)
- N = process-complexity factor = 11.5-15.5 (40 nm, 2010)

Dependability

■ Module reliability

- Mean time to failure (MTTF)
- Mean time to repair (MTTR)
- Mean time between failures (MTBF) = $MTTF + MTTR$
- Availability = $MTTF / MTBF$

Measuring Performance

- **Typical performance metrics:**
 - Response time
 - Throughput
- **Speedup of X relative to Y**
 - $\text{Execution time}_Y / \text{Execution time}_X$
- **Execution time**
 - Wall clock time: includes all system overheads
 - CPU time: only computation time
- **Benchmarks**
 - Kernels (e.g. matrix multiply)
 - Toy programs (e.g. sorting)
 - Synthetic benchmarks (e.g. Dhrystone)
 - Benchmark suites (e.g. SPEC06fp, TPC-C)

Principles of Computer Design

■ Take Advantage of Parallelism

- e.g. multiple processors, disks, memory banks, pipelining, multiple functional units

■ Principle of Locality

- Reuse of data and instructions

■ Focus on the Common Case

- Amdahl's Law

$$Execution\ time_{new} = Execution\ time_{old} \times \left((1 - Fraction_{enhanced}) + \frac{Fraction_{enhanced}}{Speedup_{enhanced}} \right)$$
$$Speedup_{overall} = \frac{1}{(1 - Fraction_{enhanced}) + \frac{Fraction_{enhanced}}{Speedup_{enhanced}}}$$

Principles of Computer Design

■ The Processor Performance Equation

CPU time = CPU clock cycles for a program × Clock cycle time

$$\text{CPU time} = \frac{\text{CPU clock cycles for a program}}{\text{Clock rate}}$$

$$\text{CPI} = \frac{\text{CPU clock cycles for a program}}{\text{Instruction count}}$$

CPU time = Instruction count × Cycles per instruction × Clock cycle time

$$\frac{\text{Instructions}}{\text{Program}} \times \frac{\text{Clock cycles}}{\text{Instruction}} \times \frac{\text{Seconds}}{\text{Clock cycle}} = \frac{\text{Seconds}}{\text{Program}} = \text{CPU time}$$

Principles of Computer Design

- Different instruction types having different CPIs

$$CPU \text{ clock cycles} = \sum_{i=1}^n IC_i \times CPI_i$$

$$CPU \text{ time} = \left(\sum_{i=1}^n IC_i \times CPI_i \right) \times \text{Clock cycle time}$$

Fallacies and Pitfalls

- **All exponential laws must come to an end**
 - Dennard scaling (constant power density)
 - Stopped by threshold voltage
 - Disk capacity
 - 30-100% per year to 5% per year
 - Moore's Law
 - Most visible with DRAM capacity
 - ITRS disbanded
 - Only four foundries left producing state-of-the-art logic chips
 - 11 nm, 3 nm might be the limit

Fallacies and Pitfalls

- **Microprocessors are a silver bullet**
 - Performance is now a programmer's burden
- **Falling prey to Amdahl's Law**
- **A single point of failure**
- **Hardware enhancements that increase performance also improve energy efficiency, or are at worst energy neutral**
- **Benchmarks remain valid indefinitely**
 - Compiler optimizations target benchmarks

Fallacies and Pitfalls

- **The rated mean time to failure of disks is 1,200,000 hours or almost 140 years, so disks practically never fail**
 - MTTF value from manufacturers assume regular replacement
- **Peak performance tracks observed performance**
- **Fault detection can lower availability**
 - Not all operations are needed for correct execution

Questions?