

Thread-level Parallelism

Professor Hugh C. Lauer

Professor Thérèse M. Smith

CS-4515, System Programming Concepts

(Slides include copyright materials from Computer Architecture: A Quantitative Approach, 5th ed., by Hennessy and Patterson and from Computer Organization and Design, 4th ed. by Patterson and Hennessy)

Flynn's Taxonomy of Multiprocessing

■ SISD

- Single instruction stream, single data stream

■ SIMD

- Single instruction stream, multiple data streams

■ MISD

- Multiple instruction streams, single data stream

■ MIMD

- Multiple instruction streams, multiple data streams

Wikipedia: Space Shuttle control system!

H&P: "No commercial system has been built ..."

Today's topic

**This Taxonomy must become part of
your Computer Science and
Engineering Vocabulary**

Questions?

What is Thread-Level Parallelism?

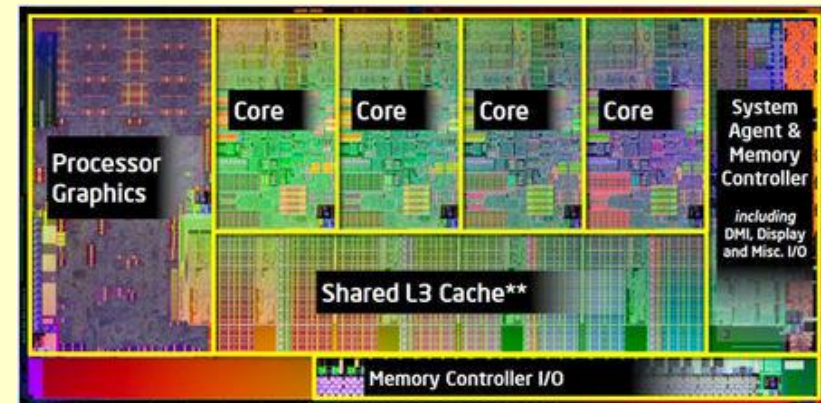
■ Multiple processors

- Not an optimization for a single processor
- Multiple Program Counters
- Shared Memory

■ MIMD model

■ Implementation

- Multithreaded processor
- Multicore processor
- Multi-processor



Why Thread-Level Parallelism?

■ Clock Speed Limit

- Plateau at 3-4 GHz in ~2005
- Thermal problems!
- Increases < 100 Mhz per year

■ ILP Limit

- Limited by interdependence of code
- Diminishing returns



Data Parallelism

- **Vector processors good for number crunching**
 - Certain kinds
- **SIMD processors**
 - Number crunching
 - Graphics
 - Imaging
 - ...
- **Imagine writing applications such as *Word*, robotic controls, web servers, airline reservation systems, or *MapReduce* in SIMD!**

Thread-level parallelism

■ SISD

- Single instruction stream, single data stream

■ SIMD

- Single instruction stream, multiple data streams

■ MISD

- Multiple instruction streams, single data stream

■ MIMD

- Multiple instruction streams, multiple data streams

Software Support

■ Requires multi-threaded software

- At least as many threads as cores

■ Easy to Parallelize

- Performing multiple independent tasks
 - Many processes on multi-purpose system
 - Handling many independent requests on server

■ Hard to Parallelize

- Intense, sequential computation
 - Won't improve completion time for one task
 - Will still improve overall throughput

Multiple threads in *MS Word* ...

- **Processing key strokes**
 - Including interpreting special characters, unicode, etc.
- **Editing the text**
- **Formatting for on-screen display**
 - Line wrap
- **Spell and grammar checking**
- **Hyphenation**
- **Pagination**
 - Widow and orphan control (!)
 - Footnotes, etc.
- ...

All in separate threads!

... and Meanwhile

- Music in the background
- E-mail and texting applications open
- Other distractions!

Process (from OS course)

- **A particular execution of a program**
 - Requires memory space
 - Processor time
 - Other resources

Thread:–

- **A particular execution of a program in the context of a process**
- **Thread vs Process in Linux**
 - ~~Thread is a special kind of process~~
 - Process is a special kind of thread

Equivalently in
Windows and MacOS

Sharing

- **(Nearly) all threads share at least some memory with some other threads!**
- **All threads share processors (and processor cores) with other threads**
- **Processor & System organization evolved to support massive multi-threaded applications**
- **Two critical issues:—**
 - Cache coherency
 - Synchronization

Questions?

Memory Model Overview

SMP

- Symmetric Multi-Processing
- **Multicore systems**
 - Typically used for $\leq 8-12$ cores
- **Physical memory (often) shared**
- **L1-L2 caches only partly shared**

DSM

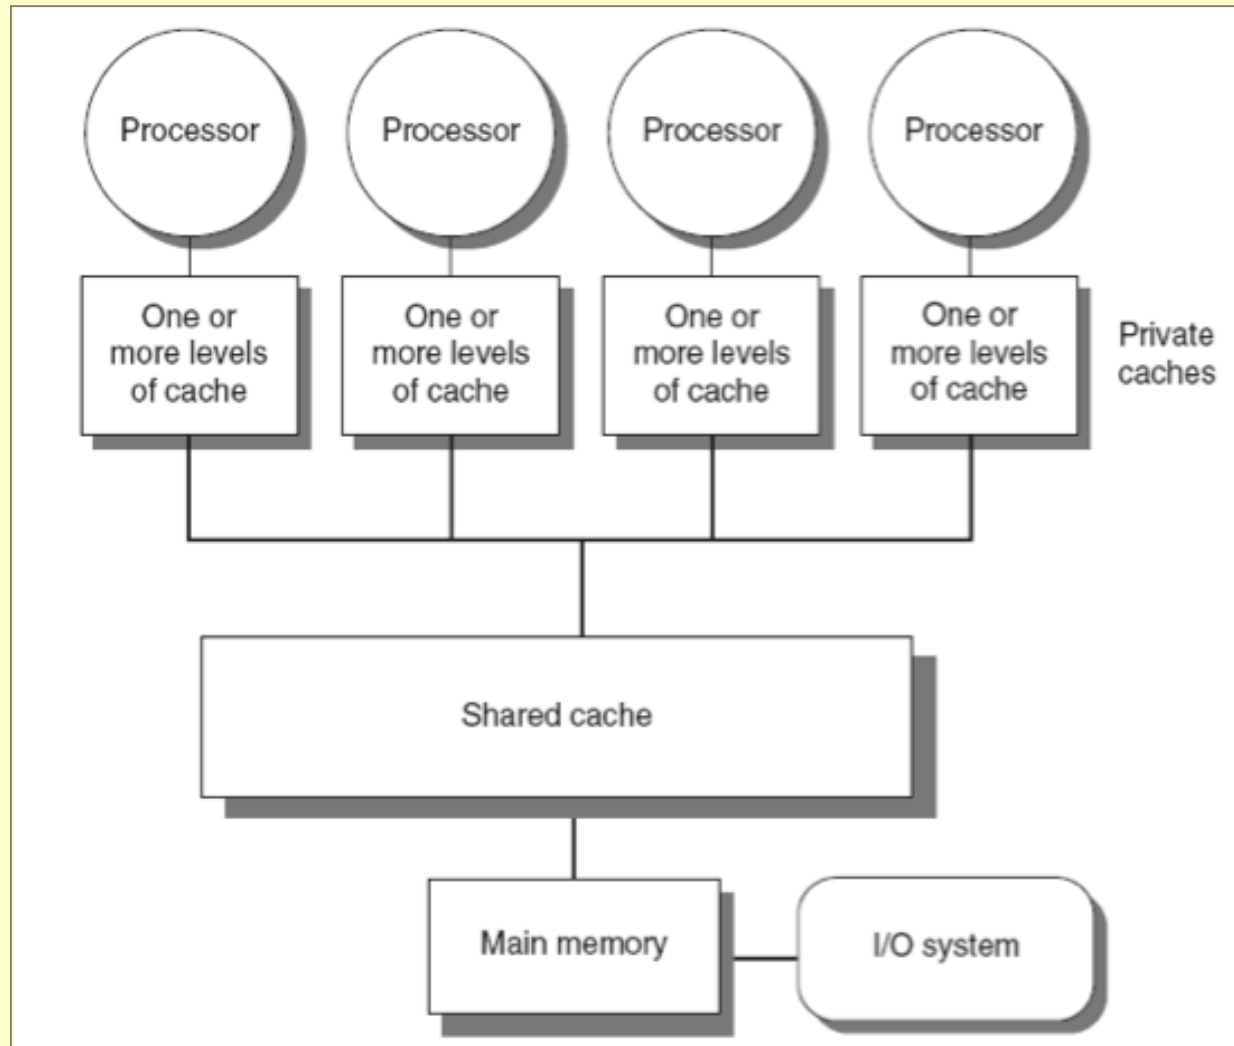
- Distributed Shared Memory
- **Multiprocessor systems**
 - Physically farther apart
- **Separate RAM per processor**
- **Bus-level links between processors**
- *Shared virtual memory*

SMP — Symmetric Multiprocessors

- **Processors have same view of shared memory**
 - Accessing memory is uniform between processors
 - Involves sharing a bus
 - Makes coherency simple

- **Good for small numbers of cores**
 - Bus bandwidth needs increase as number of CPUs grows

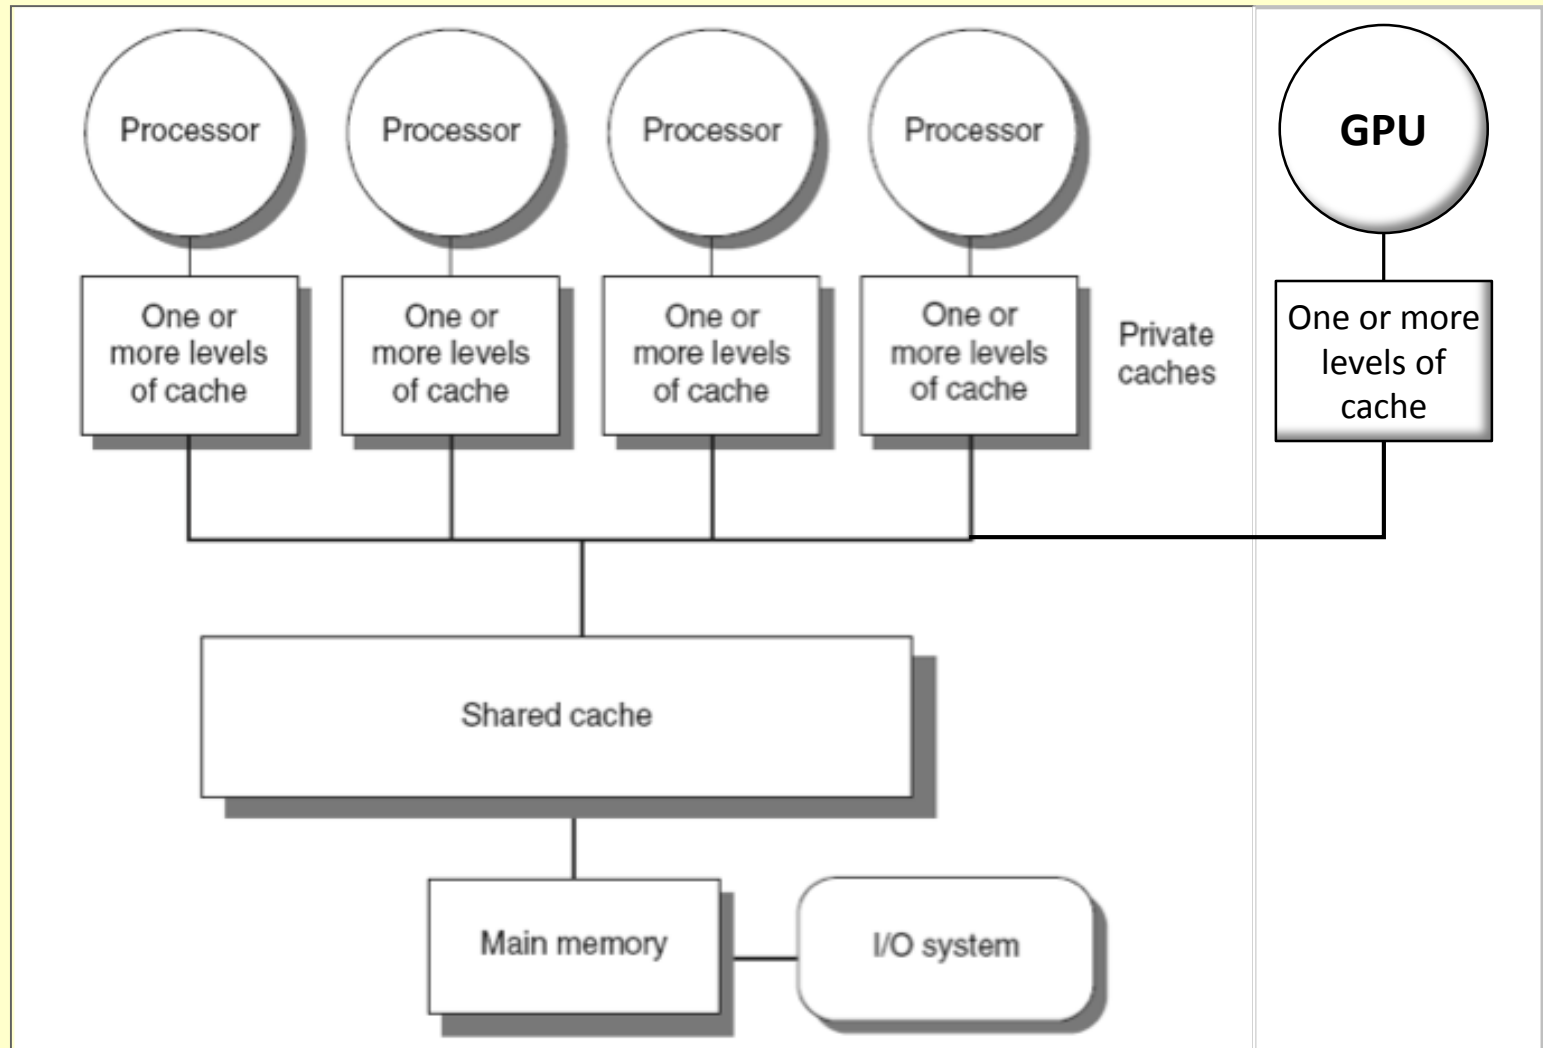
SMP



**Fig 5.1,
p. 372**

**Modern
multicore
processors**

SMP — non-uniform access patterns



DSM — Distributed Shared Memory

- **Physical memory distributed among processors**
 - Supports more processors
- **Why distribute?**
 - Memory system bandwidth
 - Reduced local latency
 - Power, cooling constraints
- **NUMA — Non-Uniform Memory Access**
 - Access time depends on memory location
 - Physically constrained by speed of light!
 - Memory is multiple cycles away ...
 - ... in each direction!

Distributed Shared Memory

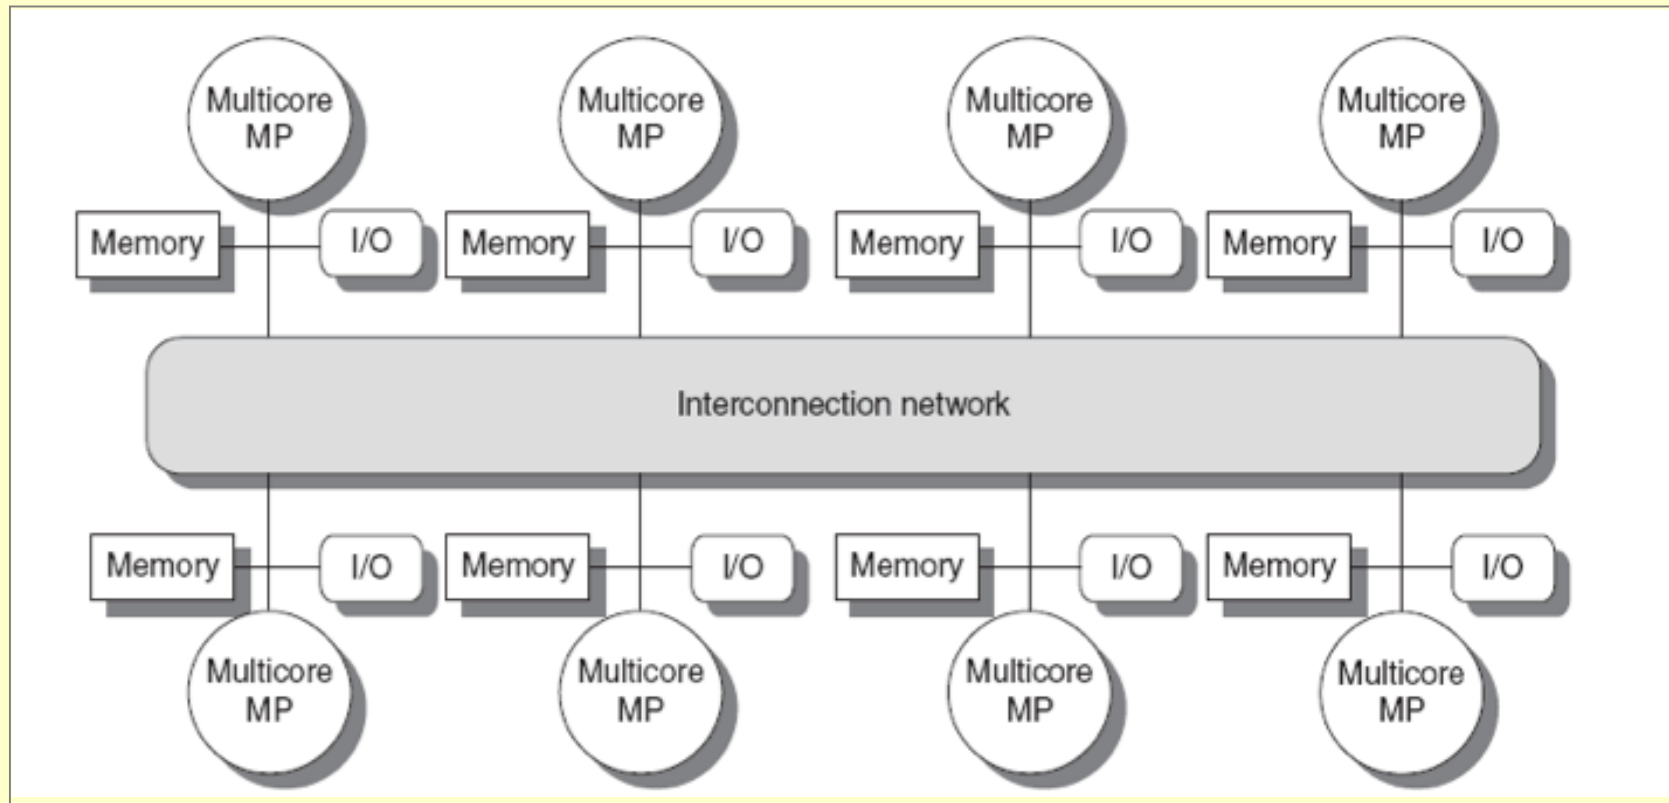


Fig. 5.2

DSM (continued)

Typical of corporate or institutional data centers

Advantages

- Memory and CPU count scalable
- Fast local memory access
- Avoids cache coherency overheads
- Cost effective

Disadvantages

- Inter-CPU communication becomes complex
 - And time-consuming!
- More effort to support in software

Questions?

Thread-level Parallelism

- Multiple processors \Rightarrow
- Multiple caches \Rightarrow
- (Possibly) Multiple main memories
- All under the control of a single operating system
- ... with many multithreaded applications

Coherency

- **What is coherency?**

- **Why do we need it?**


- **Answers:—**

- Applications are multi-threaded \Rightarrow all threads must have coherent view of application data
- Operating systems are multi-threaded \Rightarrow all threads must have coherent view of OS data structures
 - E.g., Ready queue, task lists, memory management tables, etc.

What kind of ...

- ... memory subsystem behavior do we need ...
- ... in order to support distributing a computation over multiple processors with independent memories?

Intuitive Memory Model



- Reading a memory location should return last value written to that location

- Easy in uniprocessors
 - Except for I/O

1. **Coherence** defines **values** returned by a read
2. **Consistency** determines **when** a written value will be returned by a read

1. **Coherence** defines behavior regarding same **location**
2. **Consistency** defines behavior relative to **other locations**

Defining Coherent Memory System

1. Preserve Program Order: A read by processor P to location X that follows a write by P to X, with no writes of X by another processor occurring between the write and the read by P, always returns the value written by P
2. Coherent view of memory: Read by a processor to location X that follows a write by **another processor** to X returns the written value if the read and write **are sufficiently separated in time** and no other writes to X occur between the two accesses
3. Write serialization: two writes to same location by any two processors are seen in the *same order* by all processors
 - If not, a processor could keep value last seen at write
 - For example, if the values 1 and then 2 are written to a location, processors can never read the value of the location as 2 and then later read it as 1

Digression

- **Exactly this problem can occur in, say, distributed file systems — e.g., NFS**
 - Process 1 updates a block of a file, and then signals some event seen by Process 2
 - Process 2 waits for the event, reads the block of the file, and *gets the old value* from its (file) cache

- **A fact of life with NFS!**

Write consistency

- If a processor writes location A followed by location B, ...
- ... any processor that sees the new value of B *must* also see the new value of A
- Enforced by re-order buffer
 - Writes *must* complete in program order!
 - (Reads can occur out of program order)

Questions?

Programming example

- **Thread A needs to update kernel table**

- Grabs spinlock()

- **Processor A takes an interrupt or page fault**

Time passes!

Much time passes!

- **Thread B also needs to update kernel table**

- Tries to grab spinlock
- Spins!
- ... and spins
- ... and spins
- ... and spins
- ... and spins

How does Linux kernel cope with this possibility?

Programming example 2

■ Thread A needs to add an item to application linked list

- Calls `mutex_lock()`
- Acquires lock
- Takes page fault while updating list
- Time passes
- Eventually finishes, calls `mutex_unlock()`

■ Thread B needs to add another item to same linked list

- Calls `mutex_lock()`
- Put into wait state
- Proceeds

Example 2

■ Thread A needs to add an item to application linked list

- Calls `mutex_lock()`
- Acquires lock

■ Thread B needs to add another item to same linked list

- Calls `mutex_lock()`
- Put into wait state


Two problems:—

- Calls to *`mutex_lock()`*, *`mutex_unlock()`* involve expensive calls to the operating system
 - Orders of magnitude more expensive than the data and operations that they protect
- A thread can be forced to wait for long periods of time through no fault of its own

Reality

- Highly multi-threaded applications on systems with many processors need faster, more efficient ways of synchronization
- Keep the OS out of critical paths
- Provide *wait-free* synchronization methods at application level

Example 2 (again)



Thread A

//Add new item to head of list

```

Addit: LEA  R1, newItem // &newItem
        LL   R2, head
        ST   R2, next(R1) //new -> next
        SC   R1, head
        BNEZ Addit
  
```

Thread B

//Add new item to head of list

```

Addit: LEA  R1, newItem // &newItem
        LL   R2, head
        ST   R2, next(R1) //new -> next
        SC   R1, head
        BNEZ Addit
  
```

Synchronizing operations— LL and SC


■ LL – Load Linked

- Load from memory location to a register
- Keep track of that memory location (somehow)

■ SC – Store Conditional

- Store into same memory location of most recent **LL** instruction
- Return success or failure
- *Fail* if memory location has changed since **LL**
- *Success* if *not* changed since **LL**!

Example 2 (again)



//Add new item to head of list


```
Addit: LEA  R1, newItem // &newItem
        LL   R2, head
        ST   R2, next(R1) //new -> next
        SC   R1, head
        BNEZ Addit
```

//Add new item to head of list

```
Addit: LEA  R1, newItem // &newItem
        LL   R2, head
        ST   R2, next(R1) //new -> next
        SC   R1, head
        BNEZ Addit
```

Does this work?
Did we have to invoke OS to do it?

Example 2 (more)



Thread A

//Remove 1st item from list

```

Del:  LL    R1, head // &firstItem
      LD    R2, next(R1) //head -> next
      SC    R2, head
      BNEZ Del
      // R1 contains
      &removedItem
  
```

Thread B


//Remove 1st item from list

```

Del:  LL    R1, head // & firstItem
      LD    R2, next(R1) //head -> next
      SC    R2, head
      BNEZ Del
      // R1 points to removed Item
  
```

**Does this work?
Did we have to invoke OS to do it?**

Example 2 (still more)



Thread A

//Add new item to head of list

```

Addit: LEA  R1, newItem // &newItem
        LL   R2, head
        ST   R2, next(R1) //new -> next
        SC   R1, head
        BNEZ Addit
  
```

Thread B

//Remove 1st item from list

```

Del:    LL   R1, head // & firstItem
        LD   R2, next(R1) //head -> next
        SC   R2, head
        BNEZ Del
        // R1 points to removed Item
  
```

**Does this work?
Did we have to invoke OS to do it?**

Definition

■ *Wait-free*

- Implementation of a “concurrent” object that guarantees that any process/thread can complete an operation in a finite number of steps
- No process/thread can be prevented from completing the operation by undetected halting or failures of another process
- No process/thread can be prevented from completing the operation by arbitrary variations in speed

Wait-free (continued)

■ Is Test-and-Set wait-free?

- `Test-and-Set(int *lock)`:- Atomically fetch the variable pointed to by *lock* and set it to 1 (in memory)

■ What about Fetch-and-Add?

- `Fetch-and-Add(int *lock)`:- Atomically fetch the variable pointed to by *lock* and increment it by 1 (in memory)

■ What about Exchange?

- `Exchange(int *val, int newVal)`:- Atomically fetch the variable pointed to by *val* and store the value *newVal* in its place
- See EXCH instruction, p. 415

Wait-Free (continued)

- General theory developed by Maurice Herlihy in 1991
- Compare-and-Swap is most general wait-free primitive
 - `Compare-and-Swap(*var, oldVal, newVal)`:-
Atomically fetch the variable pointed to by *var* and replace it by *newVal*, but only if its existing value equals *oldVal*
- Supports wait-free synchronization among n threads

Compare-and-swap()

- **Equivalent to *LL + SC***

- I.e., Compare-and-swap() can be implemented by LL + SC
- LL + SC can be implemented by Compare-and-swap()


- **Included in Intel architectures as *CMPEXCH***

- **Included in all IBM architectures since System 370**

- Early 1970s

- **(*LL + SC*) & *CMPEXCH* depend crucially on cache coherency mechanisms**

Example 2 (still more)



Thread A

//Add new item to head of list

```

Addit: LEA  R1, newItem // &newItem
        LL   R2, head
        ST   R2, next(R1) //new -> next
        SC   R1, head
        BNEZ Addit
  
```

Thread B

//Remove 1st item from list

```

Del:    LL   R1, head // & firstItem
        LD   R2, next(R1) //head -> next
        SC   R2, head
        BNEZ Del
        // R1 points to removed Item
  
```


Does this solution extend to n threads?

Uses of *LL + SC*

- **Atomically updating a variable by $f(\text{variable})$**
- **Atomically linking or unlinking**
- **Easy to do at the head of the list**
 - Only one item needs to be updated (i.e., head)
- **Possible to do in the middle or the tail of the list**
 - Two memory items need to be updated atomically
 - Requires clever approaches — transactions
 - ... or transactional memory (after textbook was published!)

Uses of *LL* + *SC* (continued)

- Any other management of shared data structures
- Any other management of *concurrent objects*
- Equivalent to *Compare & Swap*



Implementation of *LL* & *SC*

Questions?

Two classes of cache coherence

Snooping

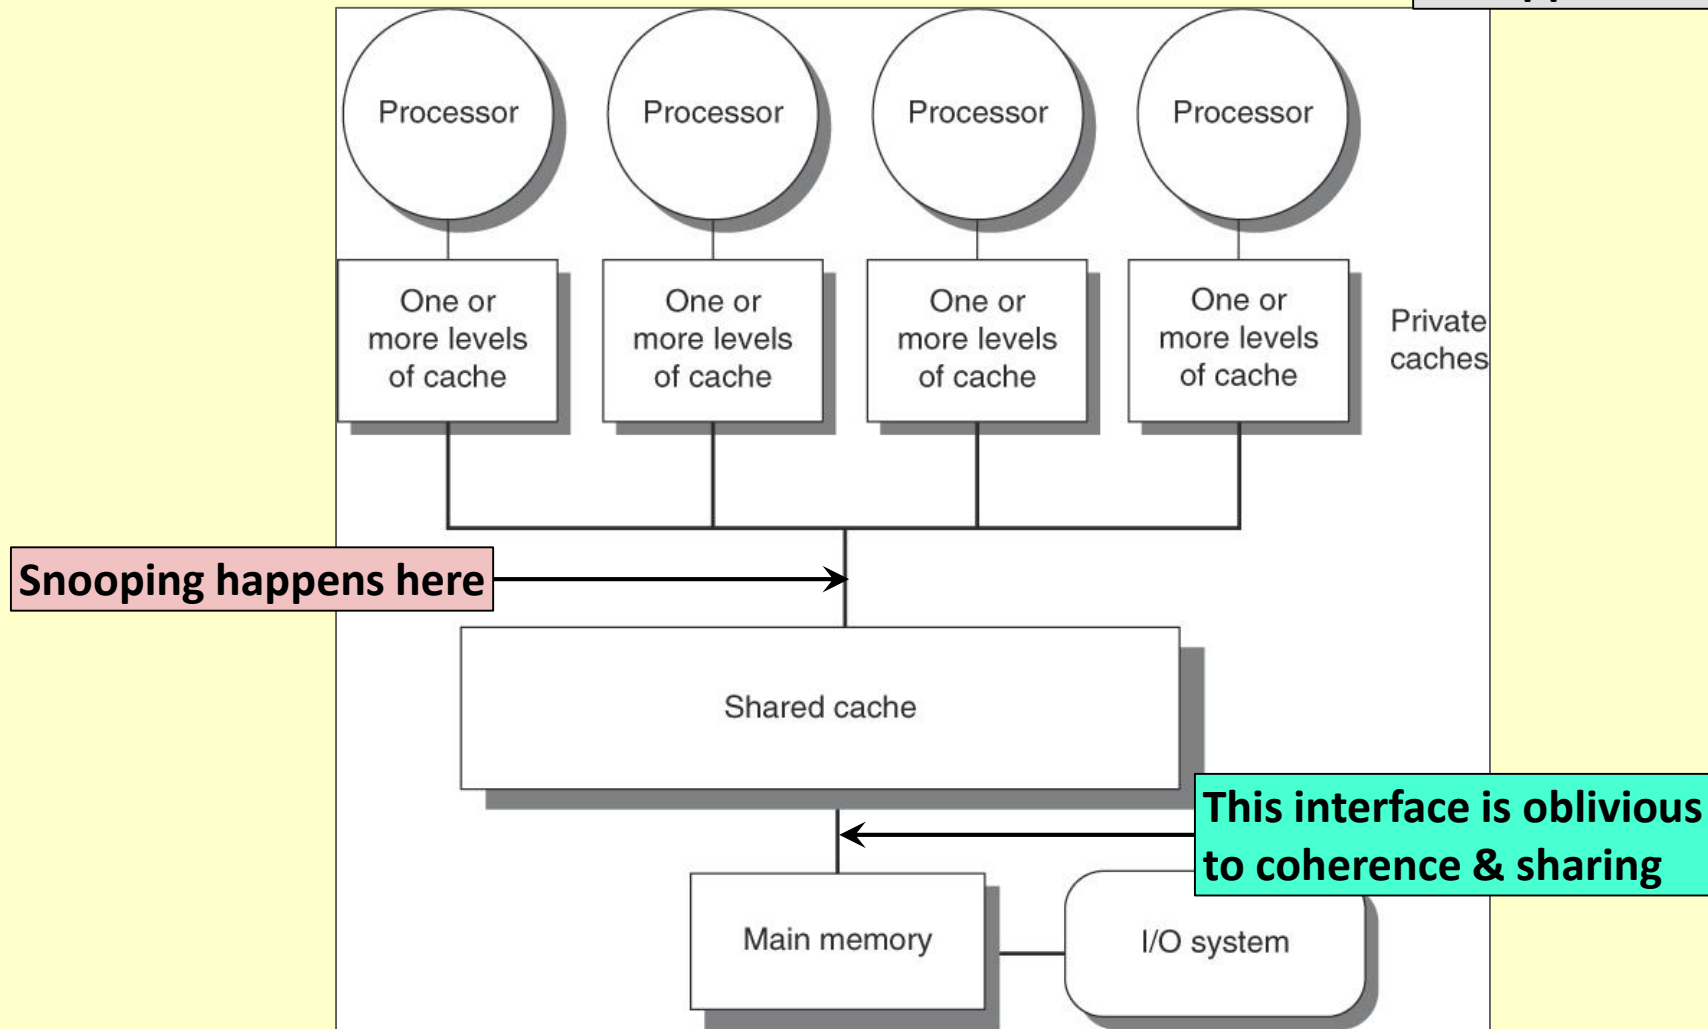
- ***Snooping*** — Every cache keeps copy of *sharing status* of each block
 - All memory access via a broadcast medium
 - Cache controllers *snoop* all transactions to update own status of cached blocks
 - *No centralized state*

Directory

- Sharing status of a block of physical memory is kept in just one location — i.e., the *directory*
 - Any processor needing to share a block *must* query its directory
 - Directory maintains info about who has *every* block

Symmetric Multi-processors

Most single-chip microprocessors use the approach



The diagram illustrates the state transitions for a 3-Cache system with three states: Invalid, Shared (read only), and Exclusive (read/write). The transitions are as follows:

- Invalid to Invalid:** CPU write hit, CPU read hit.
- Invalid to Shared:** CPU read miss, Place read miss on bus.
- Invalid to Exclusive:** CPU write miss, Place write miss on bus.
- Shared to Invalid:** Write miss for this block, Invalidate for this block.
- Shared to Shared:** CPU read hit, CPU read miss, Place read miss on bus.
- Shared to Exclusive:** CPU write miss, Place write miss on bus.
- Exclusive to Invalid:** Write-back block, Write miss for block.
- Exclusive to Shared:** CPU read miss, Write-back data; place read miss on bus.
- Exclusive to Exclusive:** CPU write hit, CPU read hit, CPU write miss, Write-back data, Place write miss on bus.

When write occurs in other caches is cannot be write ⇒ Processor gain to a data item

When write occurs, same item in other caches is invalidated — cannot be written
⇒ Processor gains exclusive access to a data item for write

Fig 5.7

Principle of write-invalidation

- **Designed for write-back caches**
- **Listens to both processor and bus**
 - Bus ensures serialization
- **Each block in cache has three possible states**
 - *Invalid* – i.e., a cache miss
 - *Shared* – block is potentially shared
 - Blocks in this state guaranteed to be up to date in memory
 - *Modified* – implies that the block is exclusive
 - Not necessarily up to date in memory
- **Assumes that all actions are atomic**
 - At the hardware level

Two components of previous diagram

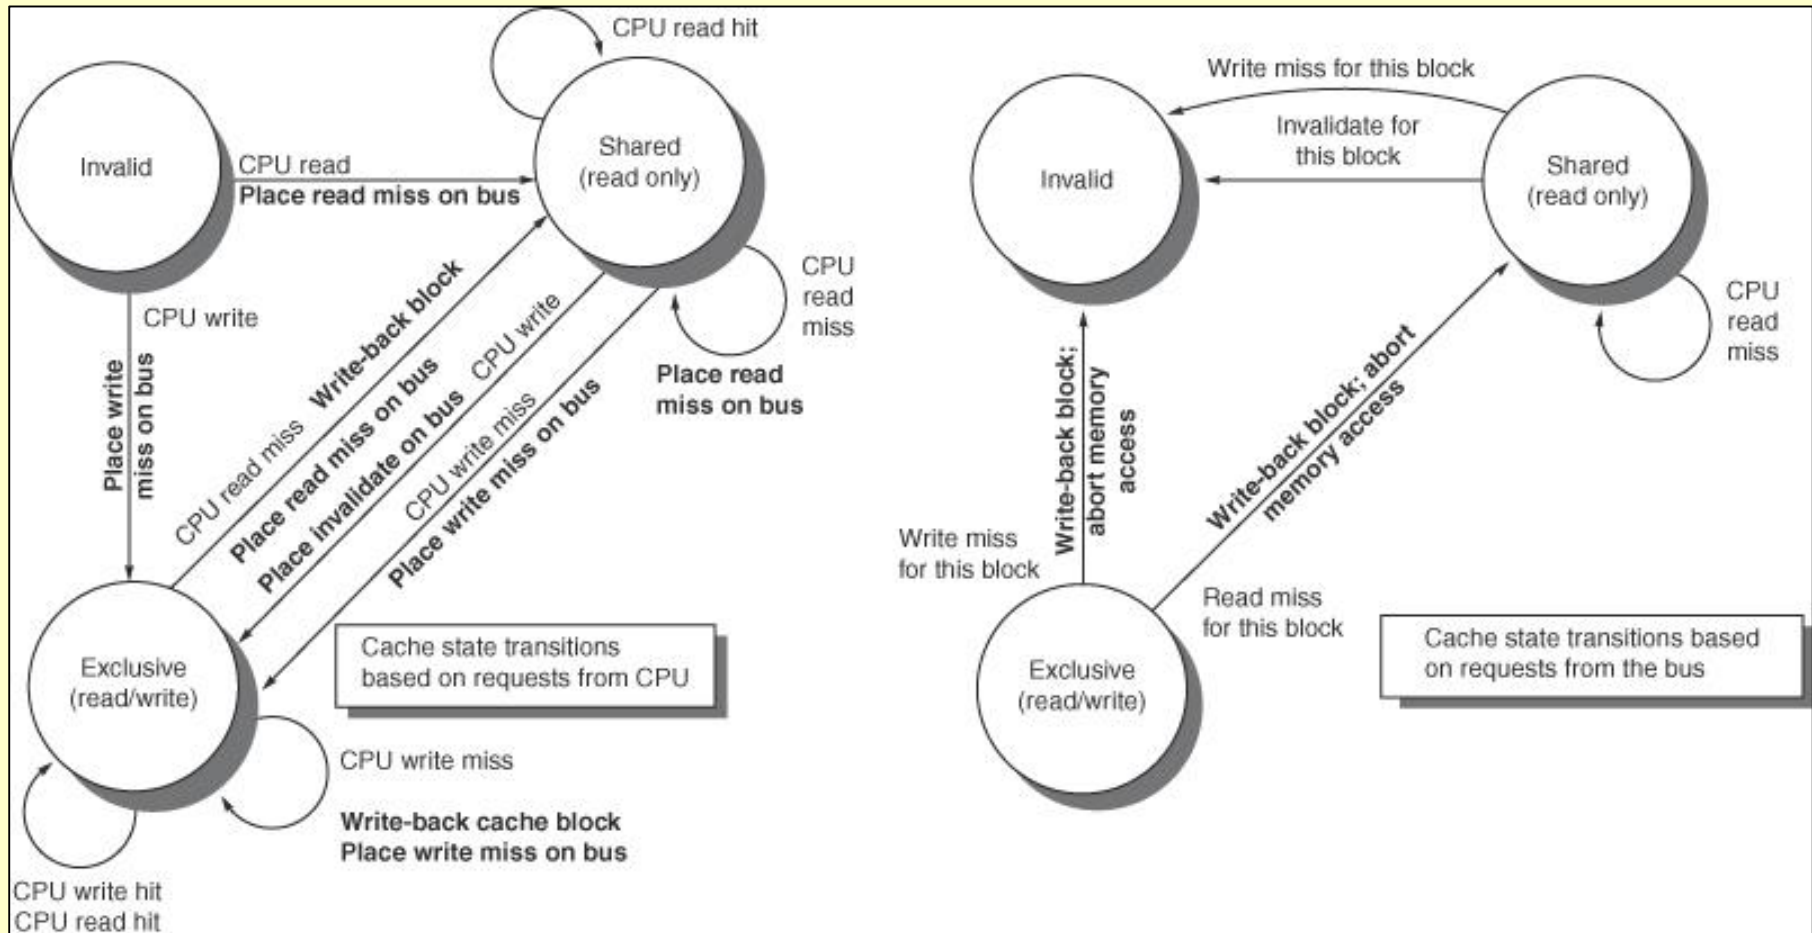


Fig 5.6

Extensions

■ MSI

- Modified (i.e., exclusive)
- Shared
- Invalid

■ MESI

- Add Exclusive but clean
 - Avoids some misses of clean data

■ MOESI

- Add owned
 - I.e., modified and dirty

Exclusivity Flag

- May be used to improve snooping
- CPU tracks whether only it has a copy of data
- If this flag is set, other snooping CPUs don't have to check their caches for copies
- Caches must now watch to see if other processors are requesting an exclusive sector
- A write-through of a block invalidates all copies and makes it exclusive again

Architectural Building Blocks

- **Cache block state transition diagram**
 - FSM specifying how disposition of block changes
 - invalid, valid, exclusive
- **Broadcast Medium Transactions (e.g., bus)**
 - Fundamental system design abstraction
 - Logically single set of wires connect several devices
 - Protocol: arbitration, command/addr, data
 - ⇒ Every device observes every transaction

□ ...

Architectural Building Blocks (continued)

- ...
- **Broadcast medium enforces serialization of read or write accesses \Rightarrow Write serialization**
 - 1st processor to get medium invalidates others copies
 - Implies cannot complete write until it obtains bus
 - All coherence schemes require serializing accesses to same cache block
- **Also need to find up-to-date copy of cache block**

The “Bus” (or other medium)

- Shared communication path among *all* processors or their caches
- Assumes an arbitration mechanism
 - Only one node can place transaction on bus at a time
 - If multiple nodes attempt at “same time,” one wins, others wait

Limitations on Busses

■ ***Bandwidth*** – a function of

- Physical size
- Lengths of wires
- Types of stubs or connections

■ ***Latency***

- Long enough to drive signal from one end to other
- Long enough to accommodate *arbitration* protocol

■ **Alternative:– *switch* with point-to-point connections**

- Like your Ethernet hub
- All nodes can see all transactions

Summary — Snooping

- Enough mechanism to implement cache coherence
- Enough mechanism to ensure memory consistency
- Underlying mechanism for implementing *LL + SC*
- Not highly scalable to many processors or multi-chip systems

Questions?

Cache Coherence and Synchronization

■ Spin lock

- If no coherence:

	DADDUI	R2,R0,#1	
lockit:	EXCH	R2,0(R1)	;atomic exchange
	BNEZ	R2,lockit	;already locked?

Spins in L1 cache
Breaks out after Invalidate

- If coherence:

lockit:	LD	R2,0(R1)	;load of lock
	BNEZ	R2,lockit	;not available-spin
	DADDUI	R2,R0,#1	;load locked value
	EXCH	R2,0(R1)	;swap
	BNEZ	R2,lockit	;branch if not 0

Cache Coherence & Synch (continued)

■ LL + SC

- Hypothetical model

Addit: LEA R1, newItem

LL R2, head

// Cache block → M or E or O

ST R2, next(R1)

SC R1, head

// Succeeds only if Cache block *still* M, E, or O!

BNEZ Addit

Implemented entirely in L1!
Depends on Cache Coherence
mechanism

Inclusion property

■ *Inclusion* \Rightarrow

- If item in L_i cache, then
must also be in L_{i+1} cache
- MSI, MESI, MOESI state must be reflected in all cache levels

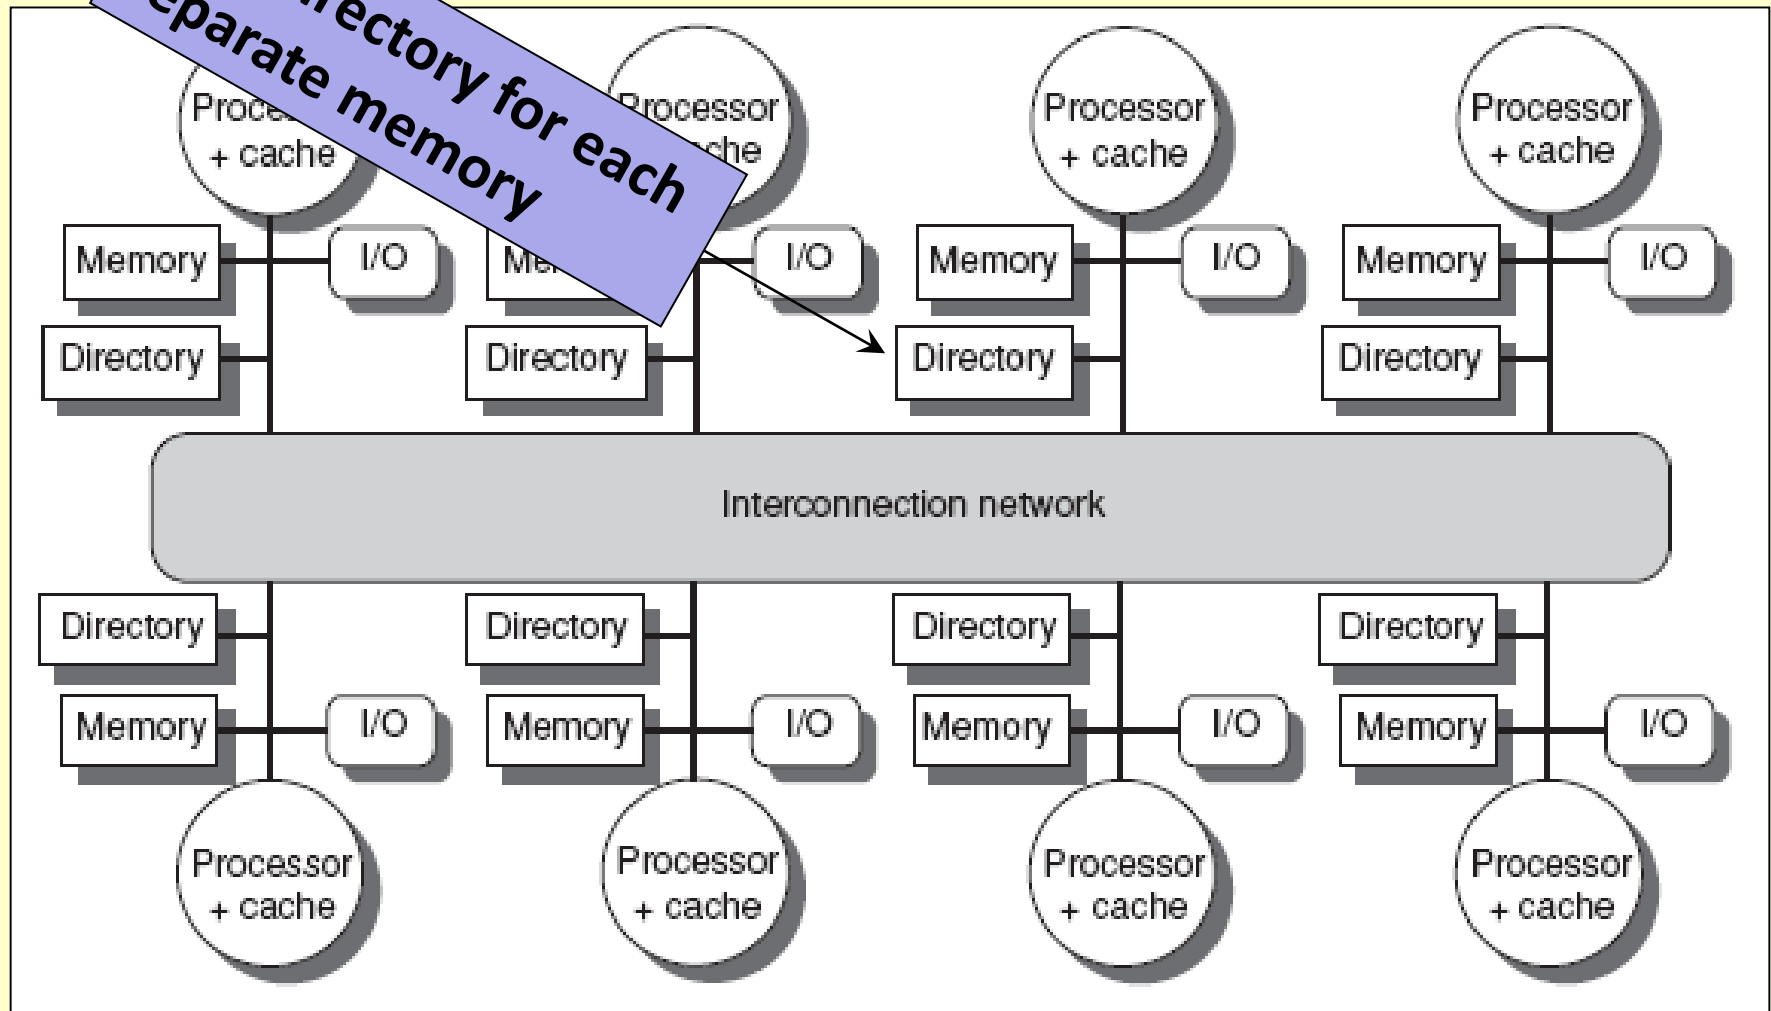
■ **LL + SC can operate locally in L1 cache if no contention**

Questions?

Scalable Approach:– Directories

- **Every memory block has associated directory information**
 - I.e., a hardware data structure
 - Tracks copies of cached blocks and their states
 - Looks up directory entry on miss, communicates only with the nodes that have copies as needed
 - In scalable networks, communication with directory and copies is via network transactions
- **Many alternatives for organizing directory information**

Directory-based Approach




<figure 5.20 p. 380>

Advantages

- **Cache misses go only to directory**
 - Not to all processors/caches
- **Directory “knows” who has what blocks**
 - Can keep status up-to-date
- **Much better use of bandwidth of interconnection network**
 - Multiple paths
 - Point-to-point networks
 - Etc.

Basic Operation of Directory



- k processors.
- With each cache-block in memory:
 k presence-bits, 1 dirty-bit
- With each cache-block in cache:
 1 valid bit, and 1 dirty (owner) bit

- Read from main memory by processor i :
 - If dirty-bit OFF then { read from main memory; turn $p[i]$ ON; }
 - if dirty-bit ON then { recall line from *owner*; update memory; turn dirty-bit OFF; turn $p[i]$ ON; supply recalled data to i ; }
- Write to main memory by processor i :
 - If dirty-bit OFF then { supply data to i ; send *invalidations* to all caches that have the block; turn dirty-bit ON; turn $p[i]$ ON; ... }

Directory Protocol

- **Similar to Snoopy Protocol: Three states**
 - *Uncached* (no processor has it; not valid in any cache)
 - *Shared*: ≥ 1 processors have data, memory up-to-date
 - *Exclusive*: 1 processor (owner) has data; memory out-of-date
- **In addition to cache state, must track which processors have data when in the shared state**
 - Usually bit vector; bit = 1 if processor has copy
- **Keep it simple(r):**
 - Writes to non-exclusive data
 \Rightarrow write miss
 - Processor blocks until access completes
 - Assume messages received and acted upon in order sent

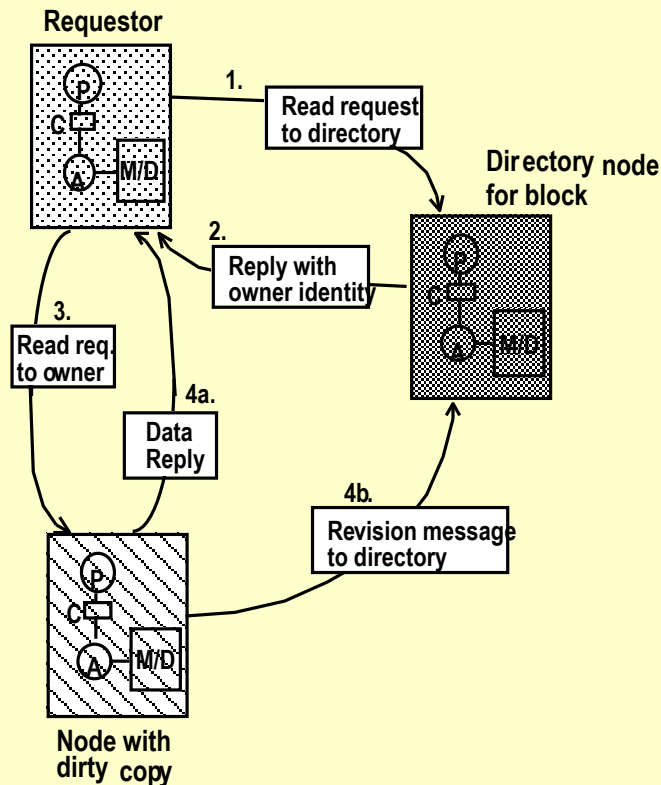
Note

- Each directory must be able to arbitrate among simultaneous, competing requests
- Requires arbitration circuitry in directory interface
 - Non-trivial

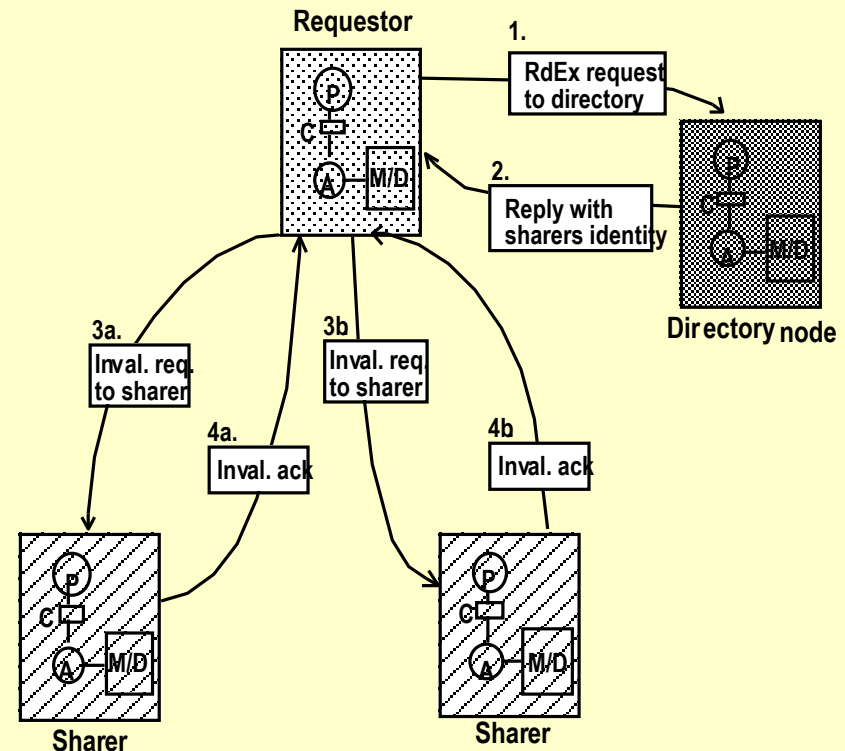
Implementing a Directory

- **Would like to assume operations atomic**
- **... but they are not;**
 - reality is much harder;
 - must avoid deadlock when running out of buffers in network (see Appendix)
- **Optimizations:—**
 - read miss or write miss in Exclusive:—
 - send data directly to requestor from owner vs. 1st to memory and then from memory to requestor

Basic Directory Transactions



(a) Read miss to a block in dirty state



(b) Write miss to a block with two sharers

See also Fig 5.22, 5.23

Questions?

Transactional Memory

- Herlihy & Moss, 1993
- Appeared in Intel & AMD processes since textbook
- More general than SC, LL

Transactional Memory (continued)

■ Like atomic transactions

- From database & distributed systems

■ *A C I D*

- *Atomicity* – to outside world, transaction happens indivisibly
- *Consistency* – transaction preserves system invariants
- *Isolated* – transactions do not interfere with each other
- *Durable* – once a transaction “commits,” the changes are permanent

Transactional Memory

■ AMD announcement of ASF

- 2006
- Not known in AMD products

■ IBM BlueGene/Q

- 2011

■ TSX instructions in Haswell


- Implementation details not published

More info to come?

Implementing LL & SC

- **LL & SC can be implemented *entirely* in L1 cache!**
 1. LL leaves cache block containing lock in “read-shared” state, *or*
 2. ... in write-exclusive state (if already in that state)
- **Any *change* to cache block by another processor or core causes state to change to *invalid* ...**
- **... thereby invalidating the LL operation**
- **Likewise, any write to same address by another thread in *same* core ...**
- **... can be detected by L1 cache to invalidate the LL operation**

Questions?



CS-4515

Computer Architecture

Professor Hugh C. Lauer
Professor Thérèse M. Smith

CS-4515, System Programming Concepts

(Slides include copyright materials from Computer Architecture: A Quantitative Approach, 6th ed., by Hennessy and Patterson and from Computer Organization and Design, 4th and 5th ed. by Patterson and Hennessy)